

STASH

#ml-papers 2019

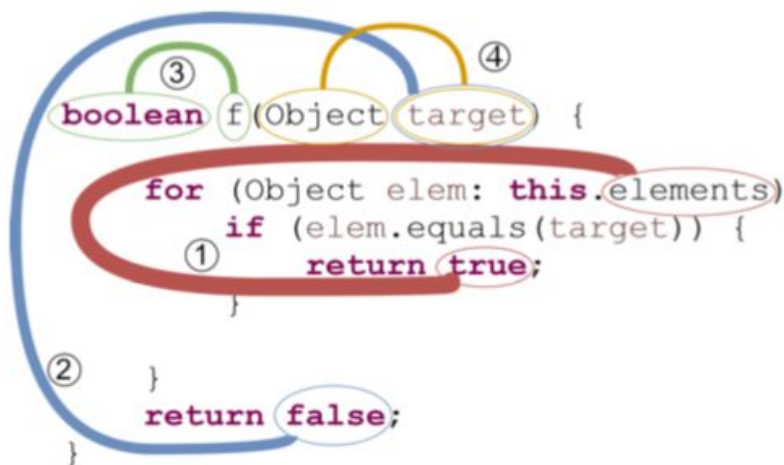
code2vec: Learning Distributed Representations of Code

Overview

- The **Code2Vec** model maps **code methods** to **vectors**, in a way that captures subtle differences.
- The algorithm works as follows:
 - a. Convert code to an **abstract syntax tree (AST)**, generated by the compiler, which reflects the codes syntax
 - b. From the AST, generate a **bag of contexts**. Each context corresponds to a path between leaf nodes.
 - c. Map each **context -> vector**
 - d. Combine all the context vectors into a single **code vector**, using a weighted sum. The weights are determined by **attention**, which is learned by the neural network.

Section 2: Motivating Example

Code Method



(a)

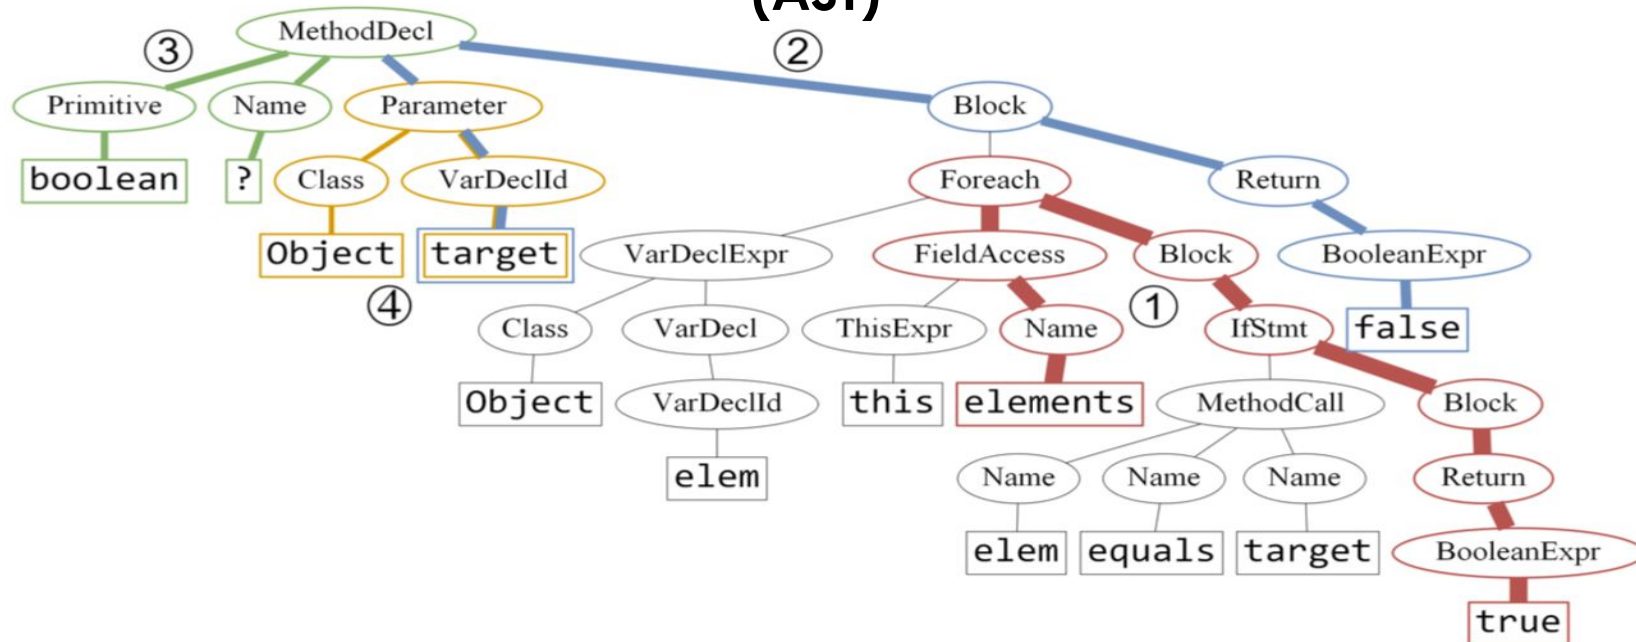
Predictions:

contains	<div><div></div></div>	90.93%
matches	<div><div></div></div>	3.54%
canHandle	<div><div></div></div>	1.15%
equals	<div><div></div></div>	0.87%
containsExact	<div><div></div></div>	0.77%

Top predictions

Abstract Syntax Tree

(AST)



- **Attention** is illustrated by the line thickness for the top contexts (1), (2), (3), (4)

Section 2: Key Ideas

- Code can be represented as a **bag of path-contexts**
- Neural network based **attention** is needed to identify the important contexts.
- The embedding vectors can be used for a lot more than predicting method names
- We can interpret the model by **visualizing the attention** of each context

Section 3: Representing Code through AST Paths


- An **abstract syntax tree (AST)** is a **compiler-generated tree** representing aspects of a program's syntax. The AST, in addition, **assigns a value to every terminal node**.
- An **AST path** of length k is a path consisting of $k+1$ nodes n_1, \dots, n_{k+1} in the graph, each connected by an edge, where the first and last nodes are **leaf** nodes, and the others are non-leaves. In addition, we include information about the direction (up, down) of each edge.
- A **path context** is defined to be a **triple**:
 - **start node value**
 - **p** (the path context)
 - **end node value**

Example 3.1. A possible path-context that represents the statement: “ $x = 7;$ ” would be:

$$\langle x, (NameExpr \uparrow AssignExpr \downarrow IntegerLiteralExpr), 7 \rangle$$

Section 4.0: Model Introduction

- **Problem:** We need to represent a code snippet (conceptually a bag of contexts) as a single vector
 - *Insight:* The should use *all* the context vectors but be allowed to learn how much focus to apply to each ([attention](#))
- **Solution:** The models performs a weighted average, i.e. performing a dot product with some global attention vector



Section 4.1-4.2: Path Contexts & Attention

- **First**, we have to turn our text code-snippets into a mathematical object
- Then, we represent C as the set of path-contexts that can be derived from it
 - Where x_s and x_t are values of AST leaves and p is the connecting path
- **Then**, we get our embedding c & on this we apply attention ($c \sim$ is post tanh activation function)
- **Finally**, we get our code vector, v as the aforementioned dot product

$$TPairs(C) = \{(term_i, term_i) \mid term_i, term_i \in termNodes(C) \wedge i \neq i\}$$

$$Rep(C) = \left\{ (x_s, p, x_t) \left| \begin{array}{l} \exists (term_s, term_t) \in TPairs(C) : \\ x_s = \phi(term_s) \wedge x_t = \phi(term_t) \\ \wedge start(p) = term_s \wedge end(p) = term_t \end{array} \right. \right\}$$

$$c_i = embedding(\langle x_s, p_j, x_t \rangle) = [value_vocab_s; path_vocab_j; value_vocab_t] \in \mathbb{R}^{3d}$$

$$\text{attention weight } \alpha_i = \frac{\exp(\tilde{c}_i^T \cdot a)}{\sum_{j=1}^n \exp(\tilde{c}_j^T \cdot a)}$$

$$\text{code vector } v = \sum_{i=1}^n \alpha_i \cdot \tilde{c}_i$$

Section 4.3-4.5: Training

- Use cross-entropy loss & gradient descent (innovative!)
- What to do with the network:
 - Create an embedding for an unseen piece of code
 - Map code into continuous space!!
- Predicting tags and names
 - way more boring
- Their **claims of contribution**:
 - Map multiple contexts into a fixed-length vector using an attention-based weighting
 - Lots of criticisms for the subtoken people (computationally inefficient, not generalizable)

Section 5: Distributed vs. Symbolic Representations

- We want to represent a qualitative concept (a word, a piece of code) quantitatively
- One approach would be to use symbolic (also referred to as localized) representations
 - each unit is uniquely represented with a single component
 - e.g one-hot-encoding
- One major downside to localized representations is high-dimensionality. For example, to represent 10-word sequences from a 100K word dictionary you need 1,000,000 dimensions
- In contrast *distributed* representations distribute the representation of an element (e.g. a word, a snippet of code) over multiple components.
 - each unit is represented by a vector describing various aspects of element
 - e.g. word embeddings
- In this model, the difference between symbolic and distributed representations of code is the difference between polynomial and linear time
 - symbolic representation $\Rightarrow O(|X|^2 * |P| * |Y|)$
 - distributed representation $\Rightarrow O(d * (|X| * |P| * |Y|))$
- Distributed representation makes training feasible!

Section 7: Model Limitations

- Closed Labels Vocabulary
 - Model is able to predict only labels that were observed as-is at training time
 - Works for the majority of targets, specifically ones which repeat across multiple programs
 - Model with struggle with more specific and diverse targets (e.g., **findUserInfoByUserIdAndKey**) and could only catch the main idea (for example: **findUserInfo**)
- Sparsity
 - Model consumes lots of trained parameters
 - large GPU memory consumption at training time
 - increases the size of the stored model
 - requires a lot of training data
 - Too much Data
 - Anything not observed in the data at training will not be represented in the model. To address this, a huge data set was used. Granularity boosts signal of less relevant data. Causes model to not perform as well as smaller datasets
- Dependency on variable names
 - When given uninformative, obfuscated or adversarial variable names, the prediction of the label is usually less accurate
 - Train the model on a mixed dataset of good and hidden variable names, hopefully reducing model dependency on variable names

Section 8: Related Work

- Bimodal modelling and natural language
 - *Bimodal*: Machines use it, humans can read it
 - They claim their way of representing the structure of code vs. other implementations treating code as an NLP problem is better
- Representation of code in ML models
 - Use syntactic relations to represent code
- Attention
 - Wide use in NLP, Speech Recognition, and image captioning
 - Attempt to make the model care about only the signal
 - E.g. white noise v. a person actually speaking
- Distributed Representations
 - Words in similar contexts have similar meanings
 - Motivated by word2vec